

# Syllabus

---

- **Introduction to SQL:** Characteristics and advantages of SQL, SQL Data Types
- **DDL Commands, DCL Commands.**
- **SQL Queries:** DML Queries with Select Query Clauses, Creating, Modifying, Deleting.
- **Views:** Creating, Dropping, Updating, Indexes,
- Set Operations, Predicates and Joins, Set membership, Grouping and Aggregation, Aggregate Functions, Nested Queries

# Characteristics of SQL

---

- SQL stands for Structured Query Language
- SQL is an ANSI and ISO standard computer language for creating and manipulating databases.
- SQL allows the user to create, update, delete, and retrieve data from a database.
- SQL is very simple and easy to learn.
- SQL works with database programs like DB2, Oracle, MS Access, Sybase, MS SQL Sever etc.
- SQL is a declarative language, not a procedural language.
- All keywords of SQL are case insensitive.

# Advantages of SQL

---

- **High Speed:** SQL Queries can be used to retrieve large amounts of records from a database quickly and efficiently.
- **Well Defined Standards Exist:** SQL databases use long-established standard, which is being adopted by ANSI & ISO. Non-SQL databases do not adhere to any clear standard.
- **No Coding Required:** Using standard SQL it is easier to manage database systems without having to write substantial amount of code.
- **Easy to learn and understand**
- **Portable:** SQL can be run on any platform, Databases using SQL can be moved from a device to another without any problems.

# SQL Data Types and Literals

**char(*n*).** Fixed length character string, with user-specified length *n*.

**varchar(*n*).** Variable length character strings, with user-specified maximum length *n*.

**Boolean.** Accepts value true or false.

**int.** Integer (a finite subset of the integers that is machine-dependent).

**smallint.** Small integer (a machine-dependent subset of the integer domain type).

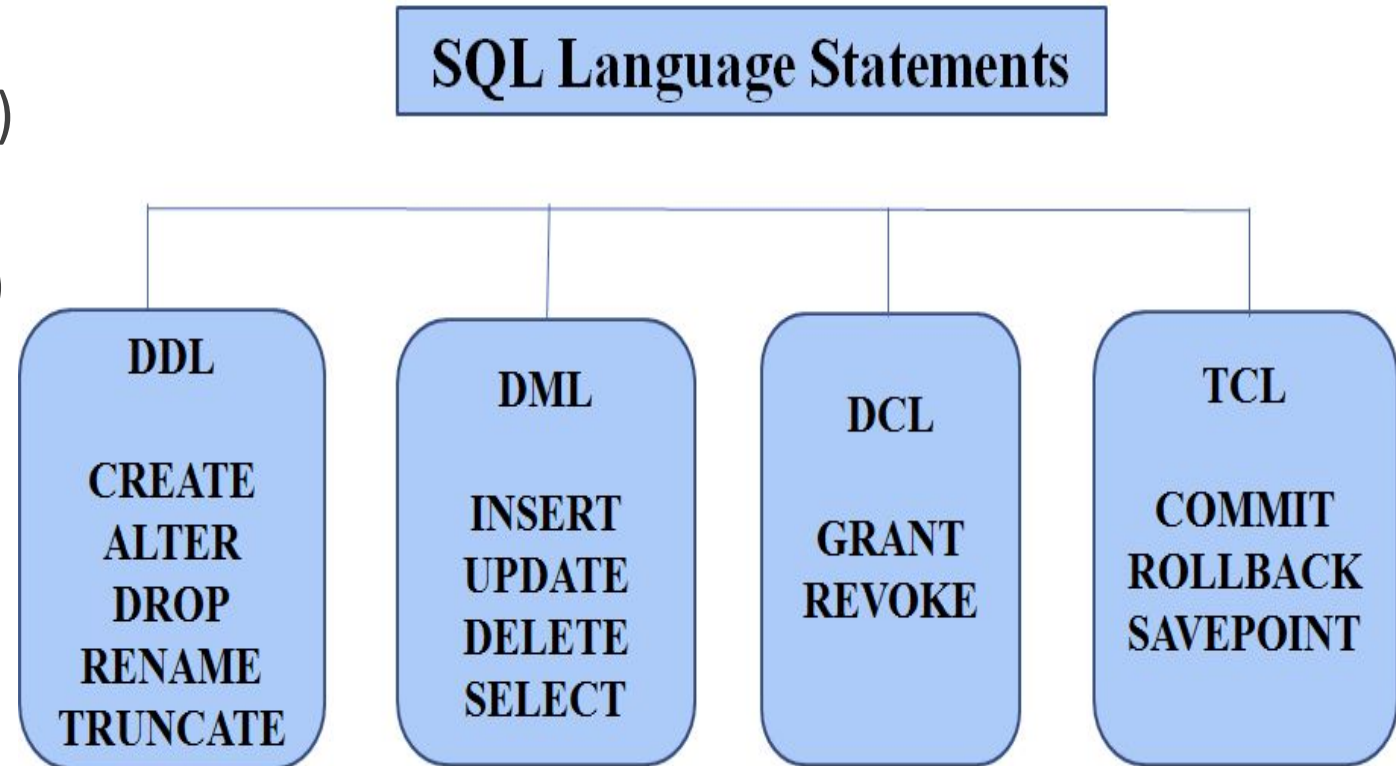
**decimal(*p,d*).** Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point. (ex., **decimal(3,1)**, allows 44.5 to be stored exactly, but not 444.5 or 0.32)

**Double(*p,d*).** Floating point and double-precision floating point numbers, with machine-dependent precision. Decimal precision can go to 53 places for a DOUBLE.

**float(*p,d*).** Floating point number, with user-specified precision of at least *n* digits. Decimal precision can go to 24 places for a FLOAT.

# SQL language statements

- Data Definition Language (DDL)
- Data Manipulation Language (DML)
- Data Control Language (DCL)
- Transaction Control Language (TCL)



# Data Definition Language (DDL)

---

- The SQL data-definition language (DDL) allows
  - Database tables to be created or deleted
  - Define indexes (keys)
  - Specify links between tables
  - Impose Integrity constraints between database tables

Some of the most commonly used DDL statements in SQL are

- **CREATE TABLE** : creates a new database tables
- **ALTER TABLE** : Alters(changes) a database tables
- **DROP TABLE** : Deletes a database table
- **RENAME TABLE** : Renames a database table
- **TRUNCATE TABLE** : Deletes all the records in the table.

# Create Table Construct

- An SQL relation is defined using the **create table** command:

```
create table  $r$  ( $A_1 D_1, A_2 D_2, \dots, A_n D_n$ ,  
    (integrity-constraint1),  
    ...,  
    (integrity-constraintk));
```


## Example:

```
create table instructor (  
    ID          char(5),  
    name        varchar(20),  
    dept_name   varchar(20),  
    salary      decimal(8,2));
```

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>

# Integrity Constraints

- Constraints are the rules enforced on the data columns of a table. These are used to limit the type of data that can go into a table. This ensures the accuracy and reliability of the data in the database.
- Constraints could be either column level or table level.
  1. **Column Level:** Column level constraints are applied to only one column.
  2. **Table Level:** Table level constraints are applied to the whole table.
- There are 4 types of Integrity Constraints:





# Domain Integrity Constraints

- Domain Integrity constraints can be defined as the definition of a valid set of values for an attribute.

- NOT NULL Constraint:**
- Unique Constraint :**
- Default Constraint :**
- Check Constraint :**

## 1. NOT NULL :

- Ensures that a column cannot have NULL value.
- E.g. Roll\_no int not null,  
Name varchar(20)

NULL value is  
not allowed

Roll_No	Name
1	ABC
2	XYZ
3	

# Domain Integrity Constraints (Cont..)

## 2. Unique Constraint :

- Ensures that all values in a column are different.
- E.g. Emp\_ID varchar(20) not null unique,

Not allowed  
as Emp\_ID  
has unique  
constraint



Emp_ID	Name	Salary
E101	ABC	20000
E102	XYZ	20000
E102	PQR	18000

## 3. Default Constraint:

- Provides a default value for a column when none is specified.
- **E.g. Marks int default NULL,**

Roll_No	Name	Marks
1	ABC	NULL
2	XYZ	NULL

# Domain Integrity Constraints (Cont..)

## 4. Check Constraint:

- The CHECK constraint ensures that all the values in a column satisfies certain conditions.

```
CREATE TABLE student (  
Roll_No int NOT NULL,  
Name varchar(255) NOT NULL,  
Age int CHECK (Age>=18)  
);
```

Roll_No	Name	Age
1	ABC	18
2	XYZ	20
3	PQR	25
4	MNP	10

**Domain Constraint**

**(Age>=18)**

**Not Allowed**

# Entity Integrity Constraints

## ■ Primary Key constraint:

- states that primary key value can't be null.
- Because primary key value is used to identify individual rows in relation and if the primary key has a null value, then we can't identify those rows.
- A table can contain a null value other than the primary key field.

Primary Key	Emp_ ID	Name	Salary
Not allowed as Emp_ID is a primary key.	E101	ABC	20000
	E102	XYZ	20000
		PQR	18000

# Referential Integrity Constraints

## Foreign Key constraint:

- A foreign key is a key used to link two tables together.
- A Foreign Key is a column or a combination of columns whose values match a Primary Key in a different table.
- The relationship between 2 tables matches the Primary Key in one of the tables with a Foreign Key in the second table.

(Table 1)

EMP_NAME	NAME	AGE	D_No
1	Jack	20	11
2	Harry	40	24
3	John	27	18
4	Devil	38	13

Foreign key

Not allowed as D\_No 18 is not defined as a Primary key of table 2 and In table 1, D\_No is a foreign key defined

Relationships

(Table 2)

Primary Key

<u>D_No</u>	D_Location
11	Mumbai
24	Delhi
13	Noida

# Referential Integrity Constraints(Cont..)

- There are two type foreign key integrity constraints:

1. cascade delete
2. cascade update

## 1. Cascade Delete :

A foreign key with cascade delete means that if a record in the parent table is deleted, then the corresponding records in the child table will automatically be deleted.

### Syntax:

```
CREATE TABLE child_table(  
    column1 datatype [ NULL | NOT NULL ],  
    column2 datatype [ NULL | NOT NULL ], ...  
    CONSTRAINT fk_name  
    FOREIGN KEY (child_col1, child_col2, ...  
    child_col_n)  
    REFERENCES parent_table (parent_col1, parent_col2,  
    ... parent_col_n)  
  
    ON DELETE CASCADE
```

```
[ ON DELETE { NO ACTION | CASCADE | SET  
NULL | SET DEFAULT } ] );
```

# Referential Integrity Constraints(Cont..)

## 2. Cascade Update :

A foreign key with cascade update means that if a record in the parent table is updated, then the corresponding records in the child table will automatically be updated.

- **DROP a FOREIGN KEY Constraint**

***ALTER TABLE ORDERS  
DROP FOREIGN KEY;***

### Syntax:

```
CREATE TABLE child_table(  
    column1 datatype [ NULL | NOT NULL ],  
    column2 datatype [ NULL | NOT NULL ], ...  
    CONSTRAINT fk_name  
    FOREIGN KEY (child_col1, child_col2, ... child_col_n)  
    REFERENCES parent_table (parent_col1, parent_col2,  
    ... parent_col_n)  
    ON UPDATE CASCADE  
    [ ON UPDATE { NO ACTION | CASCADE | SET  
    NULL | SET DEFAULT } ] );
```

# Integrity Constraints in Create Table

- **not null**
- **primary key** ( $A_1, \dots, A_n$ )
- **Foreign key** ( $A_m, \dots, A_n$ ) **references**  $r$

**Example:**

```
create table instructor (
    ID          char(5),
    name        varchar(20) not null,
    dept_name   varchar(20),
    salary      numeric(8,2),
    primary key (ID),
    foreign key (dept_name) references department);
```

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

(a) The *instructor* table

dept_name	building	budget
Comp. Sci.	Taylor	100000
Biology	Watson	90000
Elec. Eng.	Taylor	85000
Music	Packard	80000
Finance	Painter	120000
History	Painter	50000
Physics	Watson	70000

(b) The *department* table

- **primary key** declaration on an attribute automatically ensures **not null**.



# Alter Command

**Alter command is used for altering the table structure, such as,**

1. to add a column to existing table
2. to rename any existing column
3. to change datatype of any column or to modify its size.
4. to drop a column from the table.

## 2. To add a column

```
alter table <table name> add <column name> datatype
```

- All exiting tuples in the relation are assigned *null* as the value for the new attribute, if default value is not specified.
- E.g. ***ALTER TABLE Customers  
ADD Email varchar(255);***

# Alter Command (Cont..)

- By setting default value for new column

```
ALTER TABLE table_name ADD(column-name1 datatype1  
DEFAULT some_value);
```

*E.g. ALTER TABLE student ADD(  
dob DATE DEFAULT '01-Jan-99' );*

## 2.To modify a column

```
ALTER TABLE table_name modify Column(  
column_name datatype);
```

# Alter Command (Cont..)

---

E.g. *ALTER TABLE student MODIFY Column(  
address varchar(300));*

## 3. To Rename a Column

*ALTER TABLE table\_name RENAME Column  
old\_column\_name TO new\_column\_name;*

E.g. *ALTER TABLE student RENAME column  
address TO location;*

# Alter Command (Cont..)

---

## 4. To drop a column

- Dropping of attributes not supported by many databases.

```
ALTER TABLE table_name DROP Column(  
column_name);
```

- E.g. *ALTER TABLE Customers  
DROP COLUMN Email;*

# Drop Command

---

**DROP TABLE** command is used to drop an existing table in a database.

```
DROP TABLE table_name;
```

*E.g. DROP TABLE Customers;*

# Rename Command

---

**RENAME** command is used to rename a table.

```
RENAME TABLE {tbl_name} TO {new_tbl_name};
```

*E.g.    RENAME TABLE Customers TO Customers\_new;*

# Truncate Command

---

**TRUNCATE TABLE** command is used to delete complete data from an existing table.

```
TRUNCATE TABLE table_name;
```

*E.g.      TRUNCATE TABLE Customers;*

# Data Control Language (DCL)

---

**DCL commands** control the level of access that users have on database objects.

1) **GRANT** – provides access privileges to the users on the database objects. The privileges could be select, delete, update and insert on the tables and views. On the procedures, functions and packages it gives select and execute privileges.

- In DCL we have two commands,

1. **GRANT:** Used to provide any user access privileges or other privileges for the database.
2. **REVOKE:** Used to take back permissions from any user.



# GRANT Command

- **Syntax for the GRANT command :**

GRANT privilege\_name ON object\_name  
TO {user\_name | PUBLIC | role\_name} [with GRANT option];

- Allow User to create table:

To allow a user to create tables in the database, we can use the below command,

GRANT **CREATE TABLE** TO username;

- Grant Select privileges to user on customer table:

GRANT **SELECT** ON **customer** TO username;

- Grant permission to drop any table:

GRANT **DROP ANY TABLE** TO username;

- To GRANT ALL privileges to a user

GRANT **ALL PRIVILEGES** ON database\_name TO username

# DCL Example

---

```
mysql> CREATE USER 'finley'@'localhost' IDENTIFIED BY 'password';  
mysql> GRANT ALL ON *.* TO 'finley'@'localhost'; (Databasename.tablename)  
mysql> SHOW GRANTS FOR 'finley'@'localhost';  
from cmd prompt change to folder  
C:\Program Files\MySQL\MySQL Server 8.0\bin> mysql -u finley -p  
Enter password: ***** (password)  
mysql> create database a;  
mysql> use a;  
mysql> create table abc(a1 int);
```

# Continued.....

---

From Root

```
mysql> REVOKE ALL ON *.* FROM 'finley'@'localhost';
```

```
mysql> REVOKE CREATE,DROP ON *.* FROM 'finley'@'localhost';
```

# REVOKE Command

---

- **Syntax for the REVOKE command:**

```
REVOKE privilege_name ON object_name  
FROM {User_name | PUBLIC | Role_name}
```

- To take back Permissions from user

```
REVOKE CREATE TABLE FROM username;
```

- Revoke SELECT privilege on employee table from user1.

```
REVOKE SELECT ON employee FROM user1;
```

# Transaction Control Language (TCL)

---

- **TCL commands are used to** manage transactions in the database.
- These are used to manage the changes made to the data in a table by DML statements.
  - 1) Commit
  - 2) Rollback
  - 3) Savepoint

# TCL (Cont..)

---

## 1) Commit Command:

- used to permanently save any transaction into the database.
- When we use any DML command like INSERT, UPDATE or DELETE, the changes made by these commands are not permanent, until the current session is closed, the changes made by these commands can be rolled back.
- To avoid that, we use the COMMIT command to mark the changes as permanent.

- Syntax:

**COMMIT;**

# TCL (Cont..)

---

## 2) **ROLLBACK Command:**

- restores the database to last committed state.
- Can be used to cancel the last update made to the database, if those changes are not committed using the COMMIT command.
- **Syntax:**

**ROLLBACK TO savepoint\_name;**

## 3) **SAVEPOINT command:**

- used to temporarily save a transaction so that we can rollback to that point whenever required.
- **Syntax:**

**SAVEPOINT savepoint\_name;**

# Data Manipulation Language (DML)

---

**DML commands are used to make modifications of the Database like,**

- **Deletion** of tuples from a given relation.
- **Insertion** of new tuples into a given relation
- **Updation** of values in some tuples in a given relation



# INSERT Query

---

- Add a new tuple to *course*

**insert into** *course*

**values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- or equivalently

**insert into** *course* (*course\_id*, *title*, *dept\_name*, *credits*)

**values** ('CS-437', 'Database Systems', 'Comp. Sci.', 4);

- Add a new tuple to *student* with *tot\_creds* set to null

**insert into** *student*

**values** ('3003', 'Green', 'Finance', *null*);

# DELETE Query

---

- Delete all instructors from the Finance department

**delete from** *instructor*  
**where** *dept\_name*= 'Finance';

- Delete all tuples in the *student* relation.

**delete from** *student*;

# UPDATE Query

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%
- Write two **update** statements:

```
update instructor  
  set salary = salary * 0.03  
  where salary > 100000;
```

```
update instructor  
  set salary = salary * 0.05  
  where salary <= 100000;
```

- Can be done better using the **case** statement (next slide)

# Arithmetic Operations in SELECT Query

- The **select** clause can contain arithmetic expressions involving the operation,  $+$ ,  $-$ ,  $*$ , and  $/$ , and operating on constants or attributes of tuples.

- The Query:

```
select ID, name, salary/12  
from instructor;
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can rename “*salary/12*” using the **as** clause:

```
select ID, name, salary/12 as monthly_salary;
```

# SELECT Query

- The SELECT statement is used to select data from a database tables.

Select -----


From

Where


Group by

Having

Order by



Mandatory Clause



Clauses(Use as per need)

E.g. *SELECT \* FROM Student;*

- The result of an SQL query is a relation.

# SELECT Query (Cont..)

---

- An attribute can be a literal with **from** clause

*select 'A' from instructor*

- Result is a table with one column and  $N$  rows (number of tuples in the *instructors* table), each row with value “A”.

# SQL Functions

---

- Single Row Functions : Operates on each row and out for each row
  - Date Functions, String Functions, Number Functions, Conversion Functions
- Multirow Function : Aggregate Function/Group Row : Operates on Group of row, provide o/p per group

Min, Max, Count, Sum, Avg etc

- SQL Single Row Functions can be used in Select Clause, Where Clause, Group By Clause, Order By clause
- SQL Multi Row Functions can be used in Select Clause, Group By Clause, Having

# String Function : Use in Select, Where, group by , having , order by Clause

Function	Meaning
Char_length(string)	Return number of characters in argument
Concat(expr1,expr2)	Return concatenated string
Instr(expr1,expr2)	Return the index of the first occurrence of substring
Lower(expr1)	Return the argument in lowercase
Left(expr1,count)	Return the leftmost number of characters from string
Lpad(expr1,length,expr2)	left-pads a string with another string, to a certain length
Ltrim()	Remove leading spaces
Substr(string,startpos,length)	extracts a substring from a string (starting at any position).
LOCATE( <i>substring</i> , <i>string</i> , <i>start</i> )	returns the position of the first occurrence of a substring in a string
STRCMP( <i>string1</i> , <i>string2</i> )	compares two strings. Returns 0,1,-1
Upper(string)	Convert the text to upper-case
Trim(string)	removes leading and trailing spaces from a string.



# DATE Function : Use in Select, Where, group by having Clause, order by clause

Function	Meaning
ADDDATE(date, INTERVAL value addunit)	adds a time/date interval to a date and then returns the date
CURDATE() function	returns the current date. as "YYYY-MM-DD" (string)
DATEDIFF(date1, date2)	returns the number of days between two date values
DATE_SUB(date, INTERVAL value interval)	subtracts a time/date interval from a date and then returns the date
DAY(date)	returns the day of the month for a given date
DAYNAME(date)	returns the weekday name for a given date.
SYSDATE()	returns the current date and time.

# The WHERE Clause

- The **where** clause specifies conditions that the result must satisfy
  - Corresponds to the selection predicate of the relational algebra.

- To find all instructors in Comp. Sci. dept

```
select name  
from instructor  
where dept_name = 'Comp. Sci.'
```

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
  - To find all instructors in Comp. Sci. dept with salary > 80000

```
select name  
from instructor  
where dept_name = 'Comp. Sci.' and salary > 80000
```

- Comparisons can be applied to results of arithmetic expressions.

# The FROM Clause

- The **from** clause lists the relations involved in the query
  - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *instructor X teaches*

**select \***

**from** *instructor, teaches;*

- ;generates every possible instructor – teaches pair, with all attributes from both relations.
  - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)
- Cartesian product not very useful directly, but useful combined with where-clause condition (selection operation in relational algebra).

# Renaming table in Select clause

- The SQL allows renaming relations and attributes using the **as** clause:

*old-name as new-name*

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

```
select distinct T.name  
from instructor as T, instructor as S  
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```

- Keyword **as** is optional and may be omitted  
*instructor as T*  $\equiv$  *instructor T*

# Ordering the Display of Tuples

---

- List in alphabetic order the names of all instructors

```
select distinct name  
from   instructor  
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

Example: **order by** *name* **desc**

- Can sort on multiple attributes

Example: **order by** *dept\_name, name*

# Where Clause Predicates

---

- SQL includes a **between AND** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is,  $\geq$  \$90,000 and  $\leq$  \$100,000)

```
select name  
from instructor  
where salary between 90000 and 100000
```

# Null Values

- It is possible for tuples to have a null value, denoted by *null*, for some of their attributes
- *null* signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving *null* is *null*
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name  
from instructor  
where salary is null
```

# Views : Uses and Importance

- In some cases, it is not desirable for all users to see the entire logical model (i.e., all the actual relations stored in the database.)
- Consider a person who needs to know an instructors name and department, but not the salary. This person should see a relation described, in SQL, by

```
select ID, name, dept_name  
from instructor
```

- A **view** provides a mechanism to hide certain data from the view of certain users thus providing security.
- Any relation that is not of the conceptual model but is made visible to a user as a “virtual relation” is called a **view**.



# View -Syntax

- A view is defined using the **create view** statement which has the form  
**create view  $v$  as < query expression >**  
view name                      any legal SQL expression.
- Once a view is defined, the view name can be used to refer to the virtual relation that the view generates.
- Can provide huge time savings in writing queries by already having a group of frequently accessed tables joined together in a view .

# Example of Views

- A view of instructors without their salary  
**create view *faculty* as**  
**select *ID, name, dept\_name***  
**from *instructor***
- Find all instructors in the Biology department  
**select *name***  
**from *faculty***  
**where *dept\_name* = 'Biology'**
- Create a view of department salary totals  
**create view *departments\_total\_salary*(*dept\_name, total\_salary*) as**  
**select *dept\_name, sum (salary)***  
**from *instructor***  
**group by *dept\_name*;**

# Inserting a new tuple into a View

---

- Add a new tuple to *faculty* view which we defined earlier  
**insert into *faculty* values ('30765', 'Green', 'Music');**
- This insertion must be represented by the insertion of the tuple  
( '30765', 'Green', 'Music', null)  
into the *instructor* relation

# Update of a View

---

- Update query is used to Update the tuples of *a view*.

*UPDATE faculty*  
*set dept\_name="Biology"*  
*where name="ABC"*

- Updation in view reflects the original table . Means the changes will be done in the original table.

# Updatable View

---

- VIEWS are either Updatable or Read-only, but not both.
- INSERT, UPDATE, and DELETE operations are allowed on updatable VIEWS and base tables, subject to any other constraints.
- INSERT, UPDATE, and DELETE are not allowed on read-only VIEWS, but we can change their base tables, as we would expect.
- An updatable VIEW is one that can have each of its rows associated with exactly one row in an underlying base table.
- When the VIEW is changed, the changes pass unambiguously through the VIEW to that underlying base table.

# Non-Updatable View (cont..)

- In **Non-updatable view**, the SELECT statement **contain** any of the following elements:
  - Aggregate functions such as MIN, MAX, SUM, AVG, and COUNT.
  - DISTINCT
  - GROUP BY clause.
  - HAVING clause.
  - UNION or UNION ALL clause.
  - Left join or outer join.
  - Subquery in the SELECT clause or in the WHERE clause that refers to the table appeared in the FROM clause.
  - Reference to non-updatable view in the FROM clause.
  - Reference only to literal values.
  - Multiple references to any column of the base table.

# Dropping a View

---

- DROP query is used to delete a view.

## Syntax:

DROP view view\_name;

## Example:

DROP view faculty;

# Index

- Indices are data structures used to speed up access of records with specified values for index attributes.
- Indexes are used to find rows with specific column values quickly.
- Without an index, MySQL must begin with the first row and then read through the entire table to find the relevant rows. (Sequential Scan)
- If the table has an index for the columns in question, MySQL can quickly determine the position to seek to in the middle of the data file without having to look at all the data.
- This is much faster than reading every row sequentially
- MySQL create default indexes on PRIMARY KEY, UNIQUE KEY
- User defined index can be created using CREATE INDEX COMMAND
- MySQL indices are stored in B-trees.



# SQL Joins

- **Join operations** take two relations and return as a result another relation.
- A join operation is a Cartesian product which requires that tuples in the two relations match (under some condition). It also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

## *Join types*

inner join  
left outer join  
right outer join  
full outer join

## *Join Conditions*

natural  
on <predicate>  
using ( $A_1, A_1, \dots, A_n$ )

# SQL Joins : Cross Join

Cross JOIN is a simplest form of JOINS which matches each row from one database table to all rows of another. In other words it gives us combinations of each row of first table with all records in second table as cartesian product

```
SELECT * FROM `movies` CROSS JOIN `members`
```

OR

```
SELECT * FROM `movies`, `members`
```

**members**

id	first_name	last_name	movie_id
1	Adam	Smith	1
2	Ravi	Kumar	2

**movie**

movie_id	title	category
1	ASSASSIN'S CREED:	Animations
2	Real Steel(2012)	Animations

# SQL Joins : Inner Join

The inner JOIN is used to return rows from both tables that satisfy the given condition(join condition on common column ).

```
SELECT * FROM movies INNER JOIN `members` on movies.movie_id = members.movie_id
```

OR

```
SELECT * FROM movies ,members WHERE movies.movie_id = members.movie_id
```

# SQL Joins : Outer Join

MySQL Outer JOINS return all records matching from both tables .It can detect records having no match in joined table. It returns **NULL** values for records of joined table if no match is found.

```
SELECT A.title , B.first_name , B.last_name  
FROM movies "A" LEFT OUTER JOIN members "B"  
ON B.`movie_id` = A. `movie_id`
```

*# Some SQL Support keyword : Left join/natural left outer join*

**OR**

```
SELECT A.title , B.first_name , B.last_name  
FROM movies "A" LEFT OUTER JOIN members "B" USING (  
`movie_id` )
```

The LEFT JOIN returns all the rows from the table on the left even if no matching rows have been found in the table on the right.

**Where no matches have been found in the table on the right, NULL is returned.**

*What will Right Outer return?*

*What will full outer return?*


# SQL Joins



SELECT from two tables


```
SELECT *  
FROM Table1;
```

```
SELECT *  
FROM Table2;
```




INNER JOIN

```
SELECT *  
FROM Table1 t1  
INNER JOIN Table2 t2  
ON t1.fk = t2.id;
```



LEFT OUTER JOIN


```
SELECT *  
FROM Table1 t1  
LEFT OUTER JOIN Table2 t2  
ON t1.fk = t2.id;
```




RIGHT OUTER JOIN

```
SELECT *  
FROM Table1 t1  
RIGHT OUTER JOIN Table2 t2  
ON t1.fk = t2.id;
```


# SQL Joins



```
SELECT *
FROM Table1 t1
INNER JOIN Table2 t2
ON t1.fk = t2.id
INNER JOIN Table3 t3
ON t1.fk_table3 = t3.id;
```



```
SELECT *
FROM Table1 t1
LEFT OUTER JOIN Table2 t2
ON t1.fk = t2.id
LEFT OUTER JOIN Table3 t3
ON t1.fk_table3 = t3.id;
```



```
SELECT *
FROM Table1 t1
INNER JOIN Table2 t2
ON t1.fk = t2.id
LEFT OUTER JOIN Table3 t3
ON t1.fk_table3 = t3.id;
```

# Join operations – Example

Relation *course*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

Relation *prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

- Observe that  
prereq information is missing for CS-315 and  
course information is missing for CS-437



# Outer Join

---

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.



# Left Outer Join And Right Outer Join

*course* **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

• *course* **natural right outer join**

*prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

*prereq*

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

*cours*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

# Full Outer Join

- *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prereq_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

prereq			
course_id	prereq_id		
BIO-301	BIO-101		
CS-190	CS-101		
CS-347	CS-101		
cours			
course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

# Joined Relations – Examples

*Select \* from course inner join prereq on  
course.course\_id = prereq.course\_id*

course_id	title	dept_name	credits	prereq_id	course_id
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

What is the difference between the above, and a natural join?

- Select \* from course left outer join prereq on  
course.course\_id = prereq.course\_id*

course_id	title	dept_name	credits	prereq_id
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	null	null	null	CS-101

*course*

course_id	prereq_id
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

*prereq*

course_id	title	dept_name	credits
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

# Joined Relations – Examples

- *course* **full outer join** *prereq* using (*course\_id*)

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

*prereq\_id*

# Aggregate Functions

---

- Multiple –row functions are called as aggregate functions
- Act on a multiple row in the relation returning single value as an output

**avg:** average value  
**min:** minimum value  
**max:** maximum value  
**sum:** sum of values  
**count:** number of values

# Aggregate Functions Examples

Find the average salary of instructors in the Computer Science department

- **select** avg (*salary*),min(*salary*), max(*salary*),sum(*salary*)  
**from** *instructor*  
**where** *dept\_name*= 'Comp. Sci.';

Find the number of tuples in the *course* relation

- **select** count (\*) **from** *instructor*;

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000



# Aggregate Functions – Group By

Find the average salary of instructors in each department

- **select** *dept\_name*, **avg** (*salary*) **as** *avg\_salary*  
**from** *instructor*  
**group by** *dept\_name*;

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

dept_name	avg_salary
Biology	72000
Comp. Sci.	77333
Elec. Eng.	80000
Finance	85000
History	61000
Music	40000
Physics	91000

# Aggregation (Cont.)

Attributes in **select** clause outside of aggregate functions must appear in **group by** list

- /\* erroneous query \*/  
**select** *dept\_name*, *ID*, **avg** (*salary*)  
**from** *instructor*  
**group by** *dept\_name*;

Discuss why query is erroneous, [Hint :refer last table]



# Aggregate Functions – Having Clause

Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary) as avg_salary  
from instructor  
group by dept_name  
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

ID	name	dept_name	salary
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

# Like Operator

- ❑ **Percent (%):** The % character matches any substring.
- ❑ **Underscore (\_):** The character matches any character.
- ❑ Patterns are case sensitive; that is, uppercase characters do not match lowercase characters, or vice-versa.
- ❑ To illustrate pattern matching, we consider the following examples:
  - ❑ 'Intro%' matches any string beginning with “Intro”.
  - ❑ '%Comp%' matches any string containing “Comp” as a substring, for example, 'Intro. to Computer Science', and 'Computational Biology'.
  - ❑ '---' matches any string of exactly three characters.
  - ❑ '---%' matches any string of at least three characters.

# Queries with Like Operator

Using % (percent) Wildcard:

**Syntax:** `select name from table_name where column_name not like 'pattern';`

- ❖ `SELECT name FROM student WHERE city LIKE 'Luck%';`
- ❖ `SELECT name FROM student WHERE city NOT LIKE 'Luck%';`
- ❖ `SELECT * FROM student WHERE firstname = 'Ajeet' AND id > 3;`
- ❖ `SELECT * FROM student WHERE firstname = 'Ajeet' OR id > 100;`
- ❖ `SELECT * FROM student WHERE student_nm IN ('Ajeet', 'Vimal', 'Deepika');`
- ❖ `SELECT * FROM student WHERE id BETWEEN 1 AND 3;`

# Null Values and Aggregates

Total all salaries

```
select sum (salary )  
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount

All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes

- What if collection has only null values?
  - count returns 0
  - all other aggregates return null

# Subqueries (Nested Query)

A Subquery or Inner query or a Nested query is a query within another SQL query and embedded within the WHERE clause.

A subquery is used to return data that will be used in the main query as a condition to further restrict the data to be retrieved.

Subqueries can be used with the **SELECT, INSERT, UPDATE, and DELETE** statements along with the operators like =, <, >, >=, <=, IN, BETWEEN, etc.

```
SELECT ProductID,  
       Name,  
       ListPrice  
FROM   production.Product  
WHERE  ListPrice > (SELECT AVG(ListPrice)  
                   FROM   Production.Product)
```

subquery

# Examples of Subquery in DML and Select

---

SQL> SELECT \* FROM CUSTOMERS WHERE ID IN (SELECT ID FROM CUSTOMERS WHERE SALARY > 4500) ;

SQL> INSERT INTO CUSTOMERS\_BKP SELECT \* FROM CUSTOMERS WHERE ID IN (SELECT ID FROM CUSTOMERS) ;

SQL> UPDATE CUSTOMERS SET SALARY = SALARY \* 0.25 WHERE AGE IN (SELECT AGE FROM CUSTOMERS\_BKP WHERE AGE >= 27 );

SQL> DELETE FROM CUSTOMERS WHERE AGE IN (SELECT AGE FROM CUSTOMERS\_BKP WHERE AGE >= 27 );

# Subqueries in the From Clause

- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from ( select dept_name, avg (salary) as avg_salary
      from instructor
      group by dept_name)
where avg_salary > 42000;
```

Note that we do not need to use the having clause

Another way to write above query

```
select dept_name, avg_salary
from ( select dept_name, avg (salary)
      from instructor
      group by dept_name)
      as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

ID	name	dept_name	salary
22222	Einstein	Physics	95000
12121	Wu	Finance	90000
32343	El Said	History	60000
45565	Katz	Comp. Sci.	75000
98345	Kim	Elec. Eng.	80000
76766	Crick	Biology	72000
10101	Srinivasan	Comp. Sci.	65000
58583	Califieri	History	62000
83821	Brandt	Comp. Sci.	92000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
76543	Singh	Finance	80000

# Set Operations

Set operations are **union**, **intersect**, and **minus**

- Each of the above operations automatically eliminates duplicates

To retain all duplicates use the keyword all

- **union all**,
- **intersect all**
- **Minus**

table1

ID	NAME
1	ABHI
2	SAMEER
3	SAMEER

table2

- **Select name from table1 union select name from table2;**

ID	NAME
3	JAVED
4	SAMEER



# Set Operations -examples

□ SELECT name FROM table1

UNION

SELECT name FROM table2

ORDER BY name;

**Note:** The number of columns in the two selected tables or queries of a union query must match and also with similar data types.

□ SELECT id,name FROM table1

WHERE name='sameer'

UNION

SELECT id,name FROM table2

WHERE name='sameer'

ORDER BY name;

table1

ID	NAME
1	ABHI
2	SAMEER
3	SAMEER

table2

ID	NAME
3	JAVED
4	SAMEER

# Set Membership

**Find courses offered in Fall 2017 and in Spring 2018**

```
select distinct course_id  
from section  
where semester = 'Fall' and year = 2017 and  
       course_id in (select course_id  
                        from section  
                        where semester = 'Spring' and year =  
2018);
```

**Find courses offered in Fall 2017 but not in Spring 2018**

```
select distinct course_id  
from section  
where semester = 'Fall' and year = 2017 and  
       course_id not in (select course_id  
                        from section  
                        where semester = 'Spring' and year = 2018);
```

```
section (course_id , sec_id , semester ,  
         year , building , room number ,  
         time_slot_id varchar (4),  primary key  
         (course_id,  sec_id,semester,year),  
         foreign key (course id)  
         references course);
```

## Set Membership (Cont.)

- Name all instructors whose name is neither “Mozart” nor Einstein”

```
select distinct name
  from instructor
 where name not in ('Mozart', 'Einstein')
```

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>salary</i>
76766	Crick	Biology	72000
45565	Katz	Comp. Sci.	75000
10101	Srinivasan	Comp. Sci.	65000
83821	Brandt	Comp. Sci.	92000
98345	Kim	Elec. Eng.	80000
12121	Wu	Finance	90000
76543	Singh	Finance	80000
32343	El Said	History	60000
58583	Califieri	History	62000
15151	Mozart	Music	40000
33456	Gold	Physics	87000
22222	Einstein	Physics	95000

# Test for Empty Relations

---

- EXISTS and NOT EXISTS are used with a subquery in WHERE clause to examine if the result the subquery returns is TRUE or FALSE.
- The true or false value is then used to restrict the rows from outer query select.
- As EXISTS and NOT EXISTS only return TRUE or FALSE in the subquery, the SELECT list in the subquery does not need to contain actual column name(s).

## Continued...

- SELECT \* FROM customers WHERE EXISTS (SELECT \* FROM order\_details WHERE customers.customer\_id = order\_details.customer\_id);
- SELECT \* FROM customers WHERE NOT EXISTS (SELECT \* FROM order\_details WHERE customers.customer\_id = order\_details.customer\_id);
- Insert,update,delete commands can also be used with EXISTS commands
- INSERT INTO contacts (contact\_id, contact\_name) SELECT supplier\_id, supplier\_name FROM suppliers WHERE EXISTS (SELECT \* FROM orders WHERE suppliers.supplier\_id = orders.supplier\_id);
- Delete from contacts SELECT supplier\_id, supplier\_name FROM suppliers WHERE EXISTS (SELECT \* FROM orders WHERE suppliers.supplier\_id = orders.supplier\_id);

---

# The End